

Water_Potability_ML_Project

January 3, 2026

1 Water Potability Prediction System using Machine Learning

1.1 Project Information

| Item | Details |
|------------|---|
| Author | Sohaib Farooq |
| Email | sohaib.farooq@bigacademy.com |
| Domain | Environmental Science / Public Health |
| Dataset | Water Potability Dataset (Kaggle) |
| Algorithms | Logistic Regression, Random Forest, XGBoost, LightGBM |

1.2 Problem Statement

Access to safe drinking water is essential for health and is a basic human right. This project aims to **predict whether water is safe for human consumption** based on various water quality parameters using machine learning algorithms.

1.3 Dataset Download

Download Link: [Kaggle - Water Potability Dataset](#)

After downloading, place `water_potability.csv` in the same folder as this notebook.

`## Step 1 & 2: Import Required Libraries`

[43]: `# Install required packages (run this cell first time only)`
`# !pip install pandas numpy matplotlib seaborn scikit-learn`

[44]: `# Import Libraries`
`import pandas as pd`
`import numpy as np`
`import matplotlib.pyplot as plt`

```

import seaborn as sns

# Scikit-learn imports
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from sklearn.metrics import (accuracy_score, confusion_matrix,
                             classification_report, precision_score,
                             recall_score, f1_score)

# Settings
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline

print("All libraries imported successfully!")

```

All libraries imported successfully!

Step 3: Domain and Problem Statement

1.3.1 Domain: Environmental Science / Public Health

1.3.2 Problem Statement:

To predict whether water is safe for human consumption based on various water quality parameters using machine learning algorithms.

1.3.3 Features Description:

| Feature | Description | Unit |
|-------------------|---------------------------------------|------|
| pH | Acidity/Alkalinity level | 0-14 |
| Hardness | Capacity to precipitate soap | mg/L |
| Solids | Total dissolved solids | ppm |
| Chloramines | Amount of chloramines | ppm |
| Sulfate | Sulfate dissolved | mg/L |
| Conductivity | Electrical conductivity | S/cm |
| Organic_carbon | Organic carbon amount | ppm |
| Trihalomethanes | THMs amount | g/L |
| Turbidity | Cloudiness measure | NTU |
| Potability | Target: Safe (1) or Unsafe (0) | - |

Step 4: Data Collection and Preprocessing

1.3.4 4.1 Load the Dataset

```
[45]: # Load the dataset
df = pd.read_csv('water_potability.csv')

# Display basic info
print("=="*60)
print("DATASET LOADED SUCCESSFULLY!")
print("=="*60)
print(f"\nDataset Shape: {df.shape}")
print(f"Total Samples: {df.shape[0]}")
print(f"Total Features: {df.shape[1]}")
```

=====

DATASET LOADED SUCCESSFULLY!

=====

Dataset Shape: (3276, 10)
Total Samples: 3276
Total Features: 10

1.3.5 4.2 Explore the Dataset

```
[46]: # Display first 5 rows
print("First 5 Rows of Dataset:")
df.head()
```

First 5 Rows of Dataset:

```
[46]:      ph    Hardness      Solids Chloramines      Sulfate Conductivity \
0      NaN  204.890455  20791.318981     7.300212  368.516441   564.308654
1  3.716080   129.422921  18630.057858     6.635246        NaN  592.885359
2  8.099124   224.236259  19909.541732     9.275884        NaN  418.606213
3  8.316766   214.373394  22018.417441     8.059332  356.886136   363.266516
4  9.092223   181.101509  17978.986339     6.546600  310.135738   398.410813

      Organic_carbon  Trihalomethanes  Turbidity  Potability
0            10.379783       86.990970    2.963135         0
1            15.180013       56.329076    4.500656         0
2            16.868637       66.420093    3.055934         0
3            18.436524      100.341674    4.628771         0
4            11.558279       31.997993    4.075075         0
```

```
[47]: # Display last 5 rows
print("Last 5 Rows of Dataset:")
```

```
df.tail()
```

Last 5 Rows of Dataset:

```
[47]:      ph    Hardness     Solids Chloramines    Sulfate \
3271  4.668102  193.681735  47580.991603    7.166639  359.948574
3272  7.808856  193.553212  17329.802160    8.061362        NaN
3273  9.419510  175.762646  33155.578218    7.350233        NaN
3274  5.126763  230.603758  11983.869376    6.303357        NaN
3275  7.874671  195.102299  17404.177061    7.509306        NaN
```

```
      Conductivity  Organic_carbon  Trihalomethanes  Turbidity  Potability
3271      526.424171       13.894419       66.687695     4.435821         1
3272      392.449580       19.903225        NaN          2.798243         1
3273      432.044783       11.039070       69.845400     3.298875         1
3274      402.883113       11.168946       77.488213     4.708658         1
3275      327.459760       16.140368       78.698446     2.309149         1
```

```
[48]: # Dataset Information
print("Dataset Information:")
print("="*60)
df.info()
```

Dataset Information:

```
=====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3276 entries, 0 to 3275
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   ph               2785 non-null    float64
 1   Hardness          3276 non-null    float64
 2   Solids            3276 non-null    float64
 3   Chloramines       3276 non-null    float64
 4   Sulfate            2495 non-null    float64
 5   Conductivity      3276 non-null    float64
 6   Organic_carbon    3276 non-null    float64
 7   Trihalomethanes  3114 non-null    float64
 8   Turbidity          3276 non-null    float64
 9   Potability         3276 non-null    int64  
dtypes: float64(9), int64(1)
memory usage: 256.1 KB
```

```
[49]: # Statistical Summary
print("Statistical Summary:")
df.describe()
```

Statistical Summary:

```
[49]:
```

| | ph | Hardness | Solids | Chloramines | Sulfate | \ |
|-------|-------------|-------------|--------------|-------------|-------------|---|
| count | 2785.000000 | 3276.000000 | 3276.000000 | 3276.000000 | 2495.000000 | |
| mean | 7.080795 | 196.369496 | 22014.092526 | 7.122277 | 333.775777 | |
| std | 1.594320 | 32.879761 | 8768.570828 | 1.583085 | 41.416840 | |
| min | 0.000000 | 47.432000 | 320.942611 | 0.352000 | 129.000000 | |
| 25% | 6.093092 | 176.850538 | 15666.690297 | 6.127421 | 307.699498 | |
| 50% | 7.036752 | 196.967627 | 20927.833607 | 7.130299 | 333.073546 | |
| 75% | 8.062066 | 216.667456 | 27332.762127 | 8.114887 | 359.950170 | |
| max | 14.000000 | 323.124000 | 61227.196008 | 13.127000 | 481.030642 | |

| | Conductivity | Organic_carbon | Trihalomethanes | Turbidity | Potability |
|-------|--------------|----------------|-----------------|-------------|-------------|
| count | 3276.000000 | 3276.000000 | 3114.000000 | 3276.000000 | 3276.000000 |
| mean | 426.205111 | 14.284970 | 66.396293 | 3.966786 | 0.390110 |
| std | 80.824064 | 3.308162 | 16.175008 | 0.780382 | 0.487849 |
| min | 181.483754 | 2.200000 | 0.738000 | 1.450000 | 0.000000 |
| 25% | 365.734414 | 12.065801 | 55.844536 | 3.439711 | 0.000000 |
| 50% | 421.884968 | 14.218338 | 66.622485 | 3.955028 | 0.000000 |
| 75% | 481.792304 | 16.557652 | 77.337473 | 4.500320 | 1.000000 |
| max | 753.342620 | 28.300000 | 124.000000 | 6.739000 | 1.000000 |

```
[50]: # Check column names
print("Column Names:")
print(df.columns.tolist())
```

Column Names:
['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate', 'Conductivity',
'Organic_carbon', 'Trihalomethanes', 'Turbidity', 'Potability']

1.3.6 4.3 Check for Missing Values

```
[51]: # Check missing values
print("Missing Values in Each Column:")
print("-" * 60)
missing_values = df.isnull().sum()
missing_percent = (df.isnull().sum() / len(df) * 100).round(2)

missing_df = pd.DataFrame({
    'Missing Count': missing_values,
    'Percentage (%)': missing_percent
})
print(missing_df)
print(f"\nTotal Missing Values: {df.isnull().sum().sum()}")
```


Missing Values in Each Column:

| | Missing Count | Percentage (%) |
|----------|---------------|----------------|
| ph | 491 | 14.99 |
| Hardness | 0 | 0.00 |

| | | |
|-----------------|-----|-------|
| Solids | 0 | 0.00 |
| Chloramines | 0 | 0.00 |
| Sulfate | 781 | 23.84 |
| Conductivity | 0 | 0.00 |
| Organic_carbon | 0 | 0.00 |
| Trihalomethanes | 162 | 4.95 |
| Turbidity | 0 | 0.00 |
| Potability | 0 | 0.00 |

Total Missing Values: 1434

```
[52]: # Visualize missing values
plt.figure(figsize=(10, 5))
missing_values.plot(kind='bar', color='coral', edgecolor='black')
plt.title('Missing Values by Feature', fontsize=14, fontweight='bold')
plt.xlabel('Features')
plt.ylabel('Missing Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



1.3.7 4.4 Handle Missing Values

```
[53]: # Fill missing values with median (robust to outliers)
print("Handling Missing Values (Filling with Median)...")
print("="*60)

for column in df.columns:
    if df[column].isnull().sum() > 0:
```

```

median_value = df[column].median()
df[column].fillna(median_value, inplace=True)
print(f"{column}: Filled with median = {median_value:.2f}")

print("\n" + "="*60)
print("Missing Values After Treatment:")
print(df.isnull().sum())

```

Handling Missing Values (Filling with Median)...

```

ph: Filled with median = 7.04
Sulfate: Filled with median = 333.07
Trihalomethanes: Filled with median = 66.62

```

Missing Values After Treatment:

```

ph          0
Hardness    0
Solids      0
Chloramines 0
Sulfate      0
Conductivity 0
Organic_carbon 0
Trihalomethanes 0
Turbidity    0
Potability   0
dtype: int64
Sulfate: Filled with median = 333.07
Trihalomethanes: Filled with median = 66.62

```

Missing Values After Treatment:

```

ph          0
Hardness    0
Solids      0
Chloramines 0
Sulfate      0
Conductivity 0
Organic_carbon 0
Trihalomethanes 0
Turbidity    0
Potability   0
dtype: int64

```

1.3.8 4.5 Explore Target Variable

```
[54]: # Target variable distribution
print("Target Variable Distribution:")
print("-"*60)
print(df['Potability'].value_counts())
print(f"\nClass 0 (Not Potable): {(df['Potability'] == 0).sum()} samples")
print(f"Class 1 (Potable): {(df['Potability'] == 1).sum()} samples")
```

Target Variable Distribution:

```
=====
Potability
0    1998
1    1278
Name: count, dtype: int64
```


```
Class 0 (Not Potable): 1998 samples
Class 1 (Potable): 1278 samples
```

```
[55]: # Visualize target distribution
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Pie Chart
colors = ['#e74c3c', '#2ecc71']
labels = ['Not Potable (0)', 'Potable (1)']
df['Potability'].value_counts().plot(kind='pie', ax=axes[0], colors=colors,
                                      autopct='%1.1f%%', labels=labels,
                                      explode=(0.02, 0.02))
axes[0].set_title('Target Distribution (Pie Chart)', fontsize=12, fontweight='bold')
axes[0].set_ylabel('')


# Bar Chart
df['Potability'].value_counts().plot(kind='bar', ax=axes[1], color=colors, edgecolor='black')
axes[1].set_title('Target Distribution (Bar Chart)', fontsize=12, fontweight='bold')
axes[1].set_xlabel('Potability')
axes[1].set_ylabel('Count')
axes[1].set_xticklabels(labels, rotation=0)

plt.tight_layout()
plt.show()
```



1.3.9 4.6 Data Visualization

```
[56]: # Correlation Heatmap
plt.figure(figsize=(12, 8))
correlation_matrix = df.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            fmt='.2f', linewidths=0.5)
plt.title('Feature Correlation Heatmap', fontsize=14, fontweight='bold')
plt.tight_layout()
plt.show()
```




```
[57]: # Distribution of all features
fig, axes = plt.subplots(3, 3, figsize=(15, 12))
features = df.columns[:-1] # All columns except Potability

for idx, feature in enumerate(features):
    row = idx // 3
    col = idx % 3

    sns.histplot(data=df, x=feature, hue='Potability', kde=True,
                  ax=axes[row, col], palette=['#e74c3c', '#2ecc71'])
    axes[row, col].set_title(f'Distribution of {feature}')

plt.tight_layout()
plt.show()
```




```
[58]: # Box plots for all features
fig, axes = plt.subplots(3, 3, figsize=(15, 12))
features = df.columns[:-1]

for idx, feature in enumerate(features):
    row = idx // 3
    col = idx % 3

    sns.boxplot(data=df, x='Potability', y=feature, ax=axes[row, col],
                palette=['#e74c3c', '#2ecc71'])
    axes[row, col].set_title(f'{feature} by Potability')
    axes[row, col].set_xticklabels(['Not Potable', 'Potable'])

plt.tight_layout()
plt.show()
```



1.3.10 4.7 Prepare Data for Machine Learning

```
[59]: # Separate features (X) and target (y)
print("Separating Features and Target Variable...")
print("="*60)

X = df.drop('Potability', axis=1)
y = df['Potability']

print(f"Features (X) Shape: {X.shape}")
print(f"Target (y) Shape: {y.shape}")
print(f"\nFeature Names: {list(X.columns)}")
```

Separating Features and Target Variable...

=====
Features (X) Shape: (3276, 9)
Target (y) Shape: (3276,)

Feature Names: ['ph', 'Hardness', 'Solids', 'Chloramines', 'Sulfate',

```

'Conductivity', 'Organic_carbon', 'Trihalomethanes', 'Turbidity']

[60]: # Split data into training and testing sets (80% train, 20% test)
print("Splitting Data into Training and Testing Sets...")
print("*"*60)

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f"Training Set: {X_train.shape[0]} samples ({80}%)")
print(f"Testing Set: {X_test.shape[0]} samples ({20}%)")

```

Splitting Data into Training and Testing Sets...

=====

Training Set: 2620 samples (80%)
 Testing Set: 656 samples (20%)

```

[61]: # Feature Scaling using StandardScaler
print("Feature Scaling (StandardScaler)...")
print("*"*60)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Feature scaling completed successfully!")
print(f"\nScaled Training Data Shape: {X_train_scaled.shape}")
print(f"Scaled Testing Data Shape: {X_test_scaled.shape}")

```

Feature Scaling (StandardScaler)...

=====

Feature scaling completed successfully!

Scaled Training Data Shape: (2620, 9)
 Scaled Testing Data Shape: (656, 9)

Step 5: Apply ML Algorithms (4 Algorithms)

```

[62]: # Dictionary to store all results
results = {}

```

1.3.11 Algorithm 1: Logistic Regression

```

[63]: print("*"*60)
print("ALGORITHM 1: LOGISTIC REGRESSION")
print("*"*60)

```

```

# Create and train the model with class_weight='balanced' to handle class imbalance
lr_model = LogisticRegression(random_state=42, max_iter=1000, class_weight='balanced')
lr_model.fit(X_train_scaled, y_train)

# Make predictions
lr_pred = lr_model.predict(X_test_scaled)

# Calculate metrics
lr_accuracy = accuracy_score(y_test, lr_pred)
lr_precision = precision_score(y_test, lr_pred)
lr_recall = recall_score(y_test, lr_pred)
lr_f1 = f1_score(y_test, lr_pred)

# Store results
results['Logistic Regression'] = {
    'Accuracy': lr_accuracy,
    'Precision': lr_precision,
    'Recall': lr_recall,
    'F1-Score': lr_f1,
    'Predictions': lr_pred
}

# Print results
print(f"\nResults:")
print(f"    Accuracy: {lr_accuracy:.4f} ({lr_accuracy*100:.2f}%)")
print(f"    Precision: {lr_precision:.4f}")
print(f"    Recall: {lr_recall:.4f}")
print(f"    F1-Score: {lr_f1:.4f}")

```

=====

ALGORITHM 1: LOGISTIC REGRESSION

=====

Results:

```

Accuracy: 0.5259 (52.59%)
Precision: 0.4159
Recall: 0.5312
F1-Score: 0.4666

```

1.3.12 Algorithm 2: Random Forest Classifier

```
[64]: print("=="*60)
print("ALGORITHM 2: RANDOM FOREST CLASSIFIER")
print("=="*60)
```

```

# Create and train the model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_scaled, y_train)

# Make predictions
rf_pred = rf_model.predict(X_test_scaled)

# Calculate metrics
rf_accuracy = accuracy_score(y_test, rf_pred)
rf_precision = precision_score(y_test, rf_pred)
rf_recall = recall_score(y_test, rf_pred)
rf_f1 = f1_score(y_test, rf_pred)

# Store results
results['Random Forest'] = {
    'Accuracy': rf_accuracy,
    'Precision': rf_precision,
    'Recall': rf_recall,
    'F1-Score': rf_f1,
    'Predictions': rf_pred
}

# Print results
print(f"\nResults:")
print(f"  Accuracy: {rf_accuracy:.4f} ({rf_accuracy*100:.2f}%)")
print(f"  Precision: {rf_precision:.4f}")
print(f"  Recall: {rf_recall:.4f}")
print(f"  F1-Score: {rf_f1:.4f}")

```

```
=====
ALGORITHM 2: RANDOM FOREST CLASSIFIER
=====
```

Results:

```

Accuracy: 0.6585 (65.85%)
Precision: 0.6311
Recall: 0.3008
F1-Score: 0.4074

```

```
[65]: # Feature Importance from Random Forest
print("\nFeature Importance (Random Forest):")
print("="*60)

feature_importance = pd.DataFrame({
    'Feature': X.columns,
    'Importance': rf_model.feature_importances_
}).sort_values('Importance', ascending=False)
```

```


print(feature_importance.to_string(index=False))

# Visualize feature importance
plt.figure(figsize=(10, 6))
plt.barh(feature_importance['Feature'], feature_importance['Importance'], color='steelblue')
plt.xlabel('Importance')
plt.title('Feature Importance (Random Forest)', fontsize=14, fontweight='bold')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```

Feature Importance (Random Forest):

| Feature | Importance |
|-----------------|------------|
| ph | 0.129132 |
| Sulfate | 0.127858 |
| Hardness | 0.121633 |
| Chloramines | 0.116048 |
| Solids | 0.111863 |
| Conductivity | 0.100752 |
| Turbidity | 0.099522 |
| Organic_carbon | 0.097521 |
| Trihalomethanes | 0.095672 |



1.3.13 Algorithm 3: XGBoost Classifier

```
[66]: print("*"*60)
print("ALGORITHM 3: XGBOOST CLASSIFIER")
print("*"*60)

# Create and train the model
xgb_model = XGBClassifier(n_estimators=200, learning_rate=0.1, max_depth=5, random_state=42, eval_metric='logloss')
xgb_model.fit(X_train_scaled, y_train)

# Make predictions
xgb_pred = xgb_model.predict(X_test_scaled)

# Calculate metrics
xgb_accuracy = accuracy_score(y_test, xgb_pred)
xgb_precision = precision_score(y_test, xgb_pred)
xgb_recall = recall_score(y_test, xgb_pred)
xgb_f1 = f1_score(y_test, xgb_pred)

# Store results
results['XGBoost'] = {
    'Accuracy': xgb_accuracy,
    'Precision': xgb_precision,
    'Recall': xgb_recall,
    'F1-Score': xgb_f1,
    'Predictions': xgb_pred
}

# Print results
print(f"\nResults:")
print(f"  Accuracy: {xgb_accuracy:.4f} ({xgb_accuracy*100:.2f}%)")
print(f"  Precision: {xgb_precision:.4f}")
print(f"  Recall:    {xgb_recall:.4f}")
print(f"  F1-Score:  {xgb_f1:.4f}")

=====
```

ALGORITHM 3: XGBOOST CLASSIFIER

Results:

```
Accuracy: 0.6372 (63.72%)
Precision: 0.5536
Recall:    0.3633
F1-Score:  0.4387
```

1.3.14 Algorithm 4: LightGBM Classifier

```
[67]: print("=="*60)
print("ALGORITHM 4: LIGHTGBM CLASSIFIER")
print("=="*60)

# Create and train the model
lgbm_model = LGBMClassifier(n_estimators=200, learning_rate=0.1, max_depth=5, u
    ↪random_state=42, verbose=-1)
lgbm_model.fit(X_train_scaled, y_train)

# Make predictions
lgbm_pred = lgbm_model.predict(X_test_scaled)

# Calculate metrics
lgbm_accuracy = accuracy_score(y_test, lgbm_pred)
lgbm_precision = precision_score(y_test, lgbm_pred)
lgbm_recall = recall_score(y_test, lgbm_pred)
lgbm_f1 = f1_score(y_test, lgbm_pred)

# Store results
results['LightGBM'] = {
    'Accuracy': lgbm_accuracy,
    'Precision': lgbm_precision,
    'Recall': lgbm_recall,
    'F1-Score': lgbm_f1,
    'Predictions': lgbm_pred
}

# Print results
print(f"\nResults:")
print(f"  Accuracy: {lgbm_accuracy:.4f} ({lgbm_accuracy*100:.2f}%)")
print(f"  Precision: {lgbm_precision:.4f}")
print(f"  Recall:    {lgbm_recall:.4f}")
print(f"  F1-Score:  {lgbm_f1:.4f}")
```

=====

ALGORITHM 4: LIGHTGBM CLASSIFIER

=====

Results:

Accuracy: 0.6463 (64.63%)
Precision: 0.5750
Recall: 0.3594
F1-Score: 0.4423

Results:

Accuracy: 0.6463 (64.63%)

```

Precision: 0.5750
Recall:    0.3594
F1-Score:  0.4423

```

Step 6: Model Evaluation

1.3.15 6.1 Model Comparison Summary

```
[68]: # Create comparison table
print("=-*70)
print("MODEL COMPARISON SUMMARY")
print("=-*70)

algorithms = ['Logistic Regression', 'Random Forest', 'XGBoost', 'LightGBM']

comparison_df = pd.DataFrame({
    'Algorithm': algorithms,
    'Accuracy': [results[algo]['Accuracy'] for algo in algorithms],
    'Precision': [results[algo]['Precision'] for algo in algorithms],
    'Recall': [results[algo]['Recall'] for algo in algorithms],
    'F1-Score': [results[algo]['F1-Score'] for algo in algorithms]
})

# Format as percentages
comparison_df_display = comparison_df.copy()
for col in ['Accuracy', 'Precision', 'Recall', 'F1-Score']:
    comparison_df_display[col] = comparison_df_display[col].apply(lambda x: f"{x:.4f}")

print(comparison_df_display.to_string(index=False))

# Find best model
best_model_name = comparison_df.loc[comparison_df['Accuracy'].idxmax(), 'Algorithm']
best_accuracy = comparison_df['Accuracy'].max()

print("\n" + "=-*70)
print(f"BEST MODEL: {best_model_name}")
print(f"BEST ACCURACY: {best_accuracy*100:.2f}%")
print("=-*70)
```

```
=====
MODEL COMPARISON SUMMARY
=====
Algorithm Accuracy Precision Recall F1-Score
Logistic Regression 0.5259 0.4159 0.5312 0.4666
Random Forest 0.6585 0.6311 0.3008 0.4074
```

```

XGBoost    0.6372    0.5536 0.3633    0.4387
LightGBM    0.6463    0.5750 0.3594    0.4423

```

```
=====
BEST MODEL: Random Forest
BEST ACCURACY: 65.85%
=====
```

```
[69]: # Visualize model comparison
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Accuracy comparison bar chart
colors = ['#3498db', '#2ecc71', '#e74c3c', '#9b59b6']
bars = axes[0].bar(algorithms, comparison_df['Accuracy'] * 100, color=colors, ▾
    edgecolor='black')
axes[0].set_title('Model Accuracy Comparison', fontsize=14, fontweight='bold')
axes[0].set_ylabel('Accuracy (%)')
axes[0].set_ylim(0, 100)

# Add value labels on bars
for bar in bars:
    height = bar.get_height()
    axes[0].text(bar.get_x() + bar.get_width()/2., height + 1,
                 f'{height:.1f}%', ha='center', va='bottom', fontsize=10)


axes[0].tick_params(axis='x', rotation=45)

# All metrics comparison
x = np.arange(len(algorithms))
width = 0.2

axes[1].bar(x - 1.5*width, comparison_df['Accuracy'], width, label='Accuracy', ▾
    color='#3498db')
axes[1].bar(x - 0.5*width, comparison_df['Precision'], width, ▾
    label='Precision', color='#2ecc71')
axes[1].bar(x + 0.5*width, comparison_df['Recall'], width, label='Recall', ▾
    color='#e74c3c')
axes[1].bar(x + 1.5*width, comparison_df['F1-Score'], width, label='F1-Score', ▾
    color='#9b59b6')

axes[1].set_title('All Metrics Comparison', fontsize=14, fontweight='bold')
axes[1].set_ylabel('Score')
axes[1].set_xticks(x)
axes[1].set_xticklabels(algorithms, rotation=45)
axes[1].legend()
axes[1].set_ylim(0, 1)
```

```
plt.tight_layout()
plt.show()
```



1.3.16 6.2 Confusion Matrices


```
[70]: # Plot confusion matrices for all models
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

predictions = [lr_pred, rf_pred, xgb_pred, lgbm_pred]

for idx, (algo, pred) in enumerate(zip(algorithms, predictions)):
    row = idx // 2
    col = idx % 2

    cm = confusion_matrix(y_test, pred)
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=axes[row, col],
                xticklabels=['Not Potable', 'Potable'],
                yticklabels=['Not Potable', 'Potable'])
    axes[row, col].set_title(f'{algo}\nConfusion Matrix', fontsize=12,
                           fontweight='bold')
    axes[row, col].set_xlabel('Predicted')
    axes[row, col].set_ylabel('Actual')

plt.tight_layout()
plt.show()
```



1.3.17 6.3 ROC Curves and AUC Scores

```
[71]: from sklearn.metrics import roc_curve, auc, precision_recall_curve, average_precision_score

# Get prediction probabilities for ROC curves
# For models that support predict_proba
lr_proba = lr_model.predict_proba(X_test_scaled)[:, 1]
rf_proba = rf_model.predict_proba(X_test_scaled)[:, 1]
xgb_proba = xgb_model.predict_proba(X_test_scaled)[:, 1]
lgbm_proba = lgbm_model.predict_proba(X_test_scaled)[:, 1]

# Calculate ROC curves
lr_fpr, lr_tpr, _ = roc_curve(y_test, lr_proba)
rf_fpr, rf_tpr, _ = roc_curve(y_test, rf_proba)
xgb_fpr, xgb_tpr, _ = roc_curve(y_test, xgb_proba)
lgbm_fpr, lgbm_tpr, _ = roc_curve(y_test, lgbm_proba)
```

```

# Calculate AUC scores
lr_auc = auc(lr_fpr, lr_tpr)
rf_auc = auc(rf_fpr, rf_tpr)
xgb_auc = auc(xgb_fpr, xgb_tpr)
lgbm_auc = auc(lgbm_fpr, lgbm_tpr)

# Plot ROC curves
plt.figure(figsize=(10, 8))


plt.plot(lr_fpr, lr_tpr, label=f'Logistic Regression (AUC = {lr_auc:.3f})', linewidth=2)
plt.plot(rf_fpr, rf_tpr, label=f'Random Forest (AUC = {rf_auc:.3f})', linewidth=2)
plt.plot(xgb_fpr, xgb_tpr, label=f'XGBoost (AUC = {xgb_auc:.3f})', linewidth=2)
plt.plot(lgbm_fpr, lgbm_tpr, label=f'LightGBM (AUC = {lgbm_auc:.3f})', linewidth=2)

# Diagonal line (random classifier)
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier', linewidth=1)

plt.xlabel('False Positive Rate', fontsize=12)
plt.ylabel('True Positive Rate', fontsize=12)
plt.title('ROC Curves - Model Comparison', fontsize=14, fontweight='bold')
plt.legend(loc='lower right', fontsize=10)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

# Print AUC scores
print("\nAUC Scores:")
print("=="*50)
print(f"  Logistic Regression: {lr_auc:.4f}")
print(f"  Random Forest:      {rf_auc:.4f}")
print(f"  XGBoost:            {xgb_auc:.4f}")
print(f"  LightGBM:           {lgbm_auc:.4f}")

```



AUC Scores:

```
=====
Logistic Regression: 0.5475
Random Forest:      0.6407
XGBoost:            0.6367
LightGBM:           0.6487
```

1.3.18 6.4 Precision-Recall Curves

```
[72]: # Calculate Precision-Recall curves
lr_precision_curve, lr_recall_curve, _ = precision_recall_curve(y_test, ↴
    ↴lr_proba)
rf_precision_curve, rf_recall_curve, _ = precision_recall_curve(y_test, ↴
    ↴rf_proba)
xgb_precision_curve, xgb_recall_curve, _ = precision_recall_curve(y_test, ↴
    ↴xgb_proba)
lgbm_precision_curve, lgbm_recall_curve, _ = precision_recall_curve(y_test, ↴
    ↴lgbm_proba)
```

```

# Calculate Average Precision scores
lr_ap = average_precision_score(y_test, lr_proba)
rf_ap = average_precision_score(y_test, rf_proba)
xgb_ap = average_precision_score(y_test, xgb_proba)
lgbm_ap = average_precision_score(y_test, lgbm_proba)

# Plot Precision-Recall curves
plt.figure(figsize=(10, 8))


plt.plot(lr_recall_curve, lr_precision_curve, label=f'Logistic Regression (AP = {lr_ap:.3f})', linewidth=2)
plt.plot(rf_recall_curve, rf_precision_curve, label=f'Random Forest (AP = {rf_ap:.3f})', linewidth=2)
plt.plot(xgb_recall_curve, xgb_precision_curve, label=f'XGBoost (AP = {xgb_ap:.3f})', linewidth=2)
plt.plot(lgbm_recall_curve, lgbm_precision_curve, label=f'LightGBM (AP = {lgbm_ap:.3f})', linewidth=2)

# Baseline (proportion of positive class)
baseline = y_test.sum() / len(y_test)
plt.axhline(y=baseline, color='k', linestyle='--', label=f'Baseline ({baseline:.3f})', linewidth=1)

plt.xlabel('Recall', fontsize=12)
plt.ylabel('Precision', fontsize=12)
plt.title('Precision-Recall Curves - Model Comparison', fontsize=14, fontweight='bold')
plt.legend(loc='upper right', fontsize=10)
plt.grid(True, alpha=0.3)
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.tight_layout()
plt.show()

# Print Average Precision scores
print("\nAverage Precision Scores:")
print("=="*50)
print(f"  Logistic Regression: {lr_ap:.4f}")
print(f"  Random Forest:      {rf_ap:.4f}")
print(f"  XGBoost:            {xgb_ap:.4f}")
print(f"  LightGBM:           {lgbm_ap:.4f}")

```



Average Precision Scores:

```
=====
Logistic Regression: 0.4841
Random Forest:      0.5695
XGBoost:            0.5816
LightGBM:           0.5747
```

1.3.19 6.5 Classification Reports

```
[73]: # Print detailed classification reports
for algo, pred in zip(algorithms, predictions):
    print("=="*60)
    print(f"{algo} - CLASSIFICATION REPORT")
    print("=="*60)
    print(classification_report(y_test, pred,
                                target_names=['Not Potable (0)', 'Potable (1)']))
    print()
```

Logistic Regression - CLASSIFICATION REPORT

| | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Not Potable (0) | 0.64 | 0.52 | 0.57 | 400 |
| Potable (1) | 0.42 | 0.53 | 0.47 | 256 |
| accuracy | | | 0.53 | 656 |
| macro avg | 0.53 | 0.53 | 0.52 | 656 |
| weighted avg | 0.55 | 0.53 | 0.53 | 656 |

Random Forest - CLASSIFICATION REPORT

| | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Not Potable (0) | 0.66 | 0.89 | 0.76 | 400 |
| Potable (1) | 0.63 | 0.30 | 0.41 | 256 |
| accuracy | | | 0.66 | 656 |
| macro avg | 0.65 | 0.59 | 0.58 | 656 |
| weighted avg | 0.65 | 0.66 | 0.62 | 656 |

XGBoost - CLASSIFICATION REPORT

| | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Not Potable (0) | 0.67 | 0.81 | 0.73 | 400 |
| Potable (1) | 0.55 | 0.36 | 0.44 | 256 |
| accuracy | | | 0.64 | 656 |
| macro avg | 0.61 | 0.59 | 0.59 | 656 |
| weighted avg | 0.62 | 0.64 | 0.62 | 656 |

LightGBM - CLASSIFICATION REPORT

| | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Not Potable (0) | 0.67 | 0.83 | 0.74 | 400 |
| Potable (1) | 0.57 | 0.36 | 0.44 | 256 |
| accuracy | | | 0.65 | 656 |
| macro avg | 0.62 | 0.59 | 0.59 | 656 |

| | weighted avg | 0.63 | 0.65 | 0.62 | 656 |
|--|--------------|------|------|------|-----|
|--|--------------|------|------|------|-----|

Step 7: Save Results

```
[74]: # Save comparison results to CSV
comparison_df.to_csv('model_comparison_results.csv', index=False)
print("Results saved to 'model_comparison_results.csv'")
```

Results saved to 'model_comparison_results.csv'

Project Summary

```
[75]: print("*"*70)
print("                                     PROJECT SUMMARY")
print("*"*70)
print(f"""
Project Title: Water Potability Prediction System using Machine Learning

Domain: Environmental Science / Public Health

Problem Statement:
    To predict whether water is safe for human consumption based on
    various water quality parameters using machine learning algorithms.

Dataset: Water Potability Dataset (Kaggle)
    - Total Samples: {len(df)}
    - Features: {len(df.columns) - 1}
    - Target: Potability (0 = Not Potable, 1 = Potable)
```

Features Used:

- | | |
|--------------------|------------------------------------|
| 1. pH | - Acidity/Aalkalinity level (0-14) |
| 2. Hardness | - Water hardness in mg/L |
| 3. Solids | - Total dissolved solids in ppm |
| 4. Chloramines | - Chloramines amount in ppm |
| 5. Sulfate | - Sulfate dissolved in mg/L |
| 6. Conductivity | - Electrical conductivity in S/cm |
| 7. Organic_carbon | - Organic carbon in ppm |
| 8. Trihalomethanes | - THMs amount in g/L |
| 9. Turbidity | - Cloudiness measure in NTU |

Algorithms Applied:

1. Logistic Regression - Accuracy: `{results['LogisticRegression']['Accuracy'] * 100:.2f}%`

```

    2. Random Forest      - Accuracy: {results['Random Forest']['Accuracy']*100:.2f}%
    3. XGBoost            - Accuracy: {results['XGBoost']['Accuracy']*100:.2f}%
    4. LightGBM           - Accuracy: {results['LightGBM']['Accuracy']*100:.2f}%

Best Performing Model: {best_model_name}
Best Accuracy: {best_accuracy*100:.2f}%
"""
print("="*70)
print("                  PROJECT COMPLETED SUCCESSFULLY!                   ")
print("="*70)

```

=====

PROJECT SUMMARY

=====

Project Title: Water Potability Prediction System using Machine Learning

Domain: Environmental Science / Public Health

Problem Statement:

To predict whether water is safe for human consumption based on various water quality parameters using machine learning algorithms.

Dataset: Water Potability Dataset (Kaggle)

- Total Samples: 3276
- Features: 9
- Target: Potability (0 = Not Potable, 1 = Potable)

Features Used:

- | | |
|--------------------|------------------------------------|
| 1. pH | - Acidity/Aalkalinity level (0-14) |
| 2. Hardness | - Water hardness in mg/L |
| 3. Solids | - Total dissolved solids in ppm |
| 4. Chloramines | - Chloramines amount in ppm |
| 5. Sulfate | - Sulfate dissolved in mg/L |
| 6. Conductivity | - Electrical conductivity in S/cm |
| 7. Organic_carbon | - Organic carbon in ppm |
| 8. Trihalomethanes | - THMs amount in g/L |
| 9. Turbidity | - Cloudiness measure in NTU |

Algorithms Applied:

- | | |
|------------------------|--------------------|
| 1. Logistic Regression | - Accuracy: 52.59% |
| 2. Random Forest | - Accuracy: 65.85% |
| 3. XGBoost | - Accuracy: 63.72% |
| 4. LightGBM | - Accuracy: 64.63% |

Best Performing Model: Random Forest

Best Accuracy: 65.85%

=====

PROJECT COMPLETED SUCCESSFULLY!

=====

References

1. Dataset: [Kaggle - Water Potability](#)
 2. Scikit-learn Documentation: <https://scikit-learn.org/>
 3. WHO Water Quality Guidelines: <https://www.who.int/>
-

Thank You!